

# Coursera Notes for Bayesian Statistics: Techniques and Models

Jordan Katz

## Week 1

### Bayesian Modeling

data  $y$ , parameter  $\theta$

likelihood:  $p(y|\theta)$

prior:  $p(\theta)$

posterior:  $p(\theta|y)$

$$p(\theta|y) = \frac{p(y|\theta)p(\theta)}{\int p(y|\theta)p(\theta)d\theta} \propto p(y|\theta)p(\theta)$$

If we do not use conjugate priors, or if the models are more complicated, then the posterior distribution may not have a “standard” or well-known form.

### Monte Carlo Estimation

Using simulation to determine some properties of a distribution, e.g. mean, variance, probability of an event, quantiles (which all use integration)

*Example:* Suppose we have  $\theta \sim \text{Ga}(a, b)$  and want to know  $E[\theta]$

$$E[\theta] = \int_{-\infty}^{\infty} \theta p(\theta) d\theta = \int_0^{\infty} \theta \frac{b^a}{\Gamma(a)} \theta^{a-1} e^{-b\theta} d\theta = \frac{a}{b}$$

To verify with Monte Carlo, take samples  $\theta_i^*$  for  $i = 1, \dots, m$  from the Gamma distribution. Estimate sample mean as

$$\bar{\theta}^* = \frac{1}{m} \sum_{i=1}^m \theta_i^*$$

Suppose we have some function  $h(\theta)$  and we want  $E[h(\theta)]$ . Can estimate

$$E[h(\theta)] = \int h(\theta)p(\theta)d\theta \approx \frac{1}{m} \sum_{i=1}^m h(\theta_i^*)$$

In particular, if  $h(\theta)$  is  $I_A(\theta)$ , i.e. the indicator function for some event  $A$ , then we can approximate probabilities as well:  $Pr[\theta \in A]$ .

Question: How good is this estimate from sampling? By the Central Limit Theorem we know

$$\bar{\theta}^* \sim N\left(E(\theta), \frac{Var(\theta)}{m}\right)$$

The variance of the estimate is given by

$$\widehat{Var}(\theta) = \frac{1}{m} \sum_{i=1}^m (\theta_i^* - \bar{\theta}^*)^2$$

The standard error (SE) is given by

$$\sqrt{\frac{\widehat{Var}(\theta)}{m}}$$

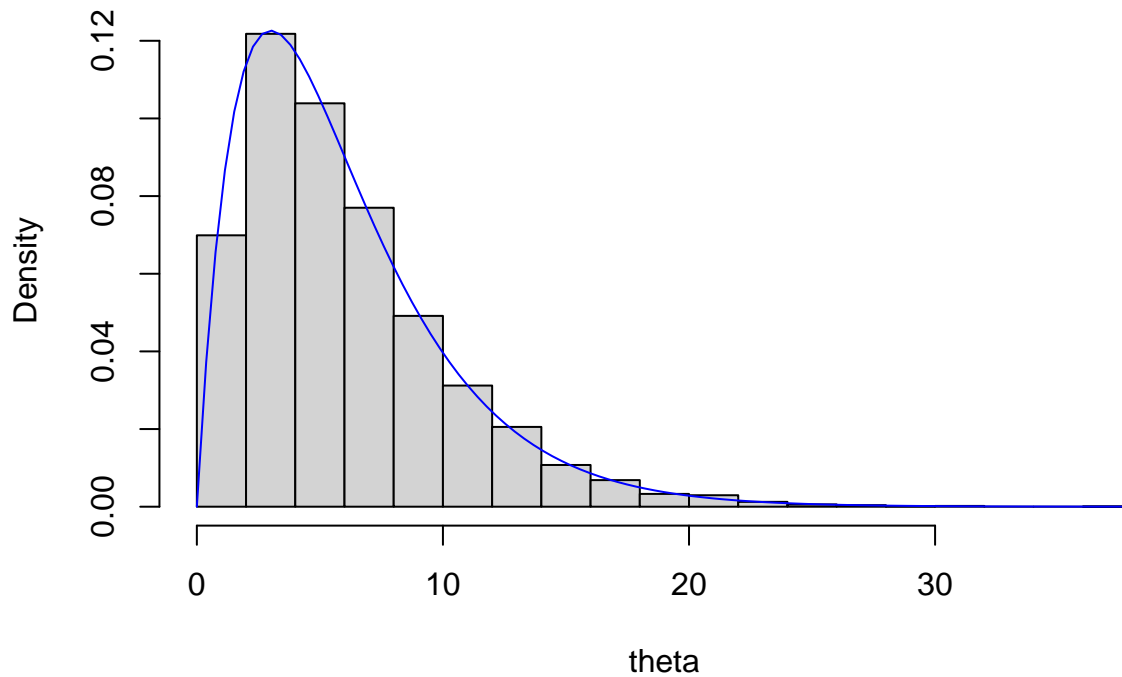
```
set.seed(32)

m=10000
a=2
b=1/3

theta = rgamma(n=m, shape=a, rate=b)

hist(theta, freq=FALSE)
curve(dgamma(x, shape=a, rate=b), col="blue", add=TRUE)
```

**Histogram of theta**



```
mean(theta) # Estimated mean
```

```
## [1] 6.022368
```

```
a/b # True mean
```

```
## [1] 6
```

```

var(theta) # Estimated variance

## [1] 18.01033
a/b^2 # True variance

## [1] 18
ind = theta < 5
mean(ind) # Estimated Prob[theta < 5]

## [1] 0.4974
pgamma(q=5, shape=a, rate=b) # True Prob[theta < 5]

## [1] 0.4963317
quantile(theta, probs=0.9) # Estimated quantile

##      90%
## 11.74426
qgamma(p=0.9, shape=a, rate=b) # True quantile

## [1] 11.66916
se = sd(theta) / sqrt(m) # Standard error of mean
mean(theta) - 2*se # Lower bound CI

## [1] 5.937491
mean(theta) + 2*se # Upper bound CI

## [1] 6.107245

```

As we can see, Monte Carlo does a pretty good job.

*Example:* Suppose we have

$$y|\phi \sim \text{Bin}(10, \phi)$$

$$\phi \sim \text{Beta}(2, 2)$$

and we want to simulate from marginal distribution of  $y$  (which can be difficult to do in general). Can do the following procedure:

1. Draw  $\phi_i^* \sim \text{Beta}(2, 2)$
2. Given  $\phi_i^*$ , draw  $y_i^* \sim \text{Bin}(10, \phi_i^*)$

Results in a list of independent pairs  $(y_i^*, \phi_i^*)$  drawn from the joint distribution. Discarding the  $\phi_i^*$ s effectively results in a sample from the marginal distribution of  $y$ .

```

m = 1e5


phi = rbeta(m, shape1=2, shape2=2)
y = rbinom(m, size=10, prob=phi)

table(y) / m

## y
##      0      1      2      3      4      5      6      7      8      9
## 0.03906 0.06925 0.09398 0.11296 0.12296 0.12412 0.12277 0.11238 0.09325 0.07005
##      10
## 0.03922

```

```
plot(table(y) / m) # Estimated marginal distribution of y
```



```
mean(y) # Estimate mean of y
```

```
## [1] 5.00046
```

## Week 2

### Metropolis-Hastings

Allows us to sample from generic distribution (whose normalizing constant may not be known). To accomplish this, we effectively construct a Markov Chain whose stationary distribution is the target distribution.

Say we want to know  $p(\theta)$  but we only know  $g(\theta)$  where  $p(\theta) \propto q(\theta)$ .

*Algorithm:*

1. Select initial value  $\theta_0$
2. for  $i = 1, \dots, m$  repeat:
  - a. Draw candidate  $\theta^* \sim q(\theta^*|\theta_{i-1})$
  - b. Define  $\alpha = \frac{g(\theta^*)/q(\theta^*|\theta_{i-1})}{g(\theta_{i-1})/q(\theta_{i-1}|\theta^*)} = \frac{g(\theta^*)}{g(\theta_{i-1})} \frac{q(\theta_{i-1}|\theta^*)}{q(\theta^*|\theta_{i-1})}$ 
    - i. if  $\alpha \geq 1$ :  
accept  $\theta^*$  and set  $\theta_i \leftarrow \theta^*$
    - ii.  $0 < \alpha < 1$ :  
with prob  $\alpha$ : accept  $\theta^*$  and set  $\theta_i \leftarrow \theta^*$   
with prob  $1 - \alpha$ : reject  $\theta^*$  and set  $\theta_i \leftarrow \theta_{i-1}$

Where  $q$  here is the candidate generating distribution which may or may not depend on  $\theta_{i-1}$ .

One choice is to make  $q$  the same distribution regardless of the value  $\theta_{i-1}$ . If we take this option, we want  $q(\theta)$  to be similar to  $p(\theta)$  to best approximate it. A high acceptance rate is a good sign here but still may want  $q$  to have a larger variance than  $p$  to assure we are exploring the space well.

Another choice – one which *does* depend on  $\theta_{i-1}$  – is to choose a distribution  $q$  that is centered on  $\theta_{i-1}$ .

A common choice for such a distribution is  $N(\theta_{i-1}, 1)$ , or in other words, a Gaussian random walk:  $\theta^* = \theta_{i-1} + N(0, 1)$  In this particular case, we have

$$q(\theta^*|\theta_{i-1}) = \frac{1}{\sqrt{2\pi}} \exp[-0.5(\theta^* - \theta_{i-1})^2] = q(\theta_{i-1}|\theta^*)$$

The “size” of the random walk step can affect acceptance (and thus convergence) rate. A high acceptance rate is not a good sign here. If random walk is taking too small of steps, it will accept candidate more often but will take a long time to fully explore the space. If it is taking too large of steps, many proposals will have low probabilities which leads to a low acceptance rate. This amounts to “wasted” samples. Ideally, a random walk sampler should have an acceptance rate between 23% and 50%.

In any symmetric case, we have the property  $q(a|b) = q(b|a)$ , so step 2 in the algorithm above reduces to

$$\alpha = \frac{g(\theta^*)}{g(\theta_{i-1})}$$